

```

[hadoop@ip-172-31-5-137 wordcount]$ hadoop jar WordCount2.jar
com.wordcount2.WordCount s3://biblewordcount/ s3://biblewordcount/output/
18/09/30 06:19:55 INFO client.RMProxy: Connecting to ResourceManager at
ip-172-31-5-137.us-east-2.compute.internal/172.31.5.137:8032
18/09/30 06:19:56 WARN mapreduce.JobResourceUploader: Hadoop command-line
option parsing not performed. Implement the Tool interface and execute your
application with ToolRunner to remedy this.
18/09/30 06:19:57 INFO input.FileInputFormat: Total input files to process :
11
18/09/30 06:19:57 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library
18/09/30 06:19:57 INFO lzo.LzoCodec: Successfully loaded & initialized
native-lzo library [hadoop-lzo rev 29dbe82f1ecfd8384b89f1a32e5e2d9e5dfd7724]
18/09/30 06:19:57 INFO mapreduce.JobSubmitter: number of splits:11
18/09/30 06:19:57 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1538278733852_0006
18/09/30 06:19:57 INFO impl.YarnClientImpl: Submitted application
application_1538278733852_0006
18/09/30 06:19:57 INFO mapreduce.Job: The url to track the job: http://
ip-172-31-5-137.us-east-2.compute.internal:20888/proxy/
application_1538278733852_0006/
18/09/30 06:19:57 INFO mapreduce.Job: Running job: job_1538278733852_0006
^C[hadoop@ip-172-31-5-137 wordcount]$ hadoop jar WordCount2.jar
com.wordcount2.WordCount s3://biblewordcount/ s3://biblewordcount/tput/
18/09/30 06:20:45 INFO client.RMProxy: Connecting to ResourceManager at
ip-172-31-5-137.us-east-2.compute.internal/172.31.5.137:8032
18/09/30 06:20:46 WARN mapreduce.JobResourceUploader: Hadoop command-line
option parsing not performed. Implement the Tool interface and execute your
application with ToolRunner to remedy this.
18/09/30 06:20:46 INFO input.FileInputFormat: Total input files to process :
9
18/09/30 06:20:46 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library
18/09/30 06:20:46 INFO lzo.LzoCodec: Successfully loaded & initialized
native-lzo library [hadoop-lzo rev 29dbe82f1ecfd8384b89f1a32e5e2d9e5dfd7724]
18/09/30 06:20:46 WARN hdfs.DataStreamer: Caught exception
java.lang.InterruptedExecution
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1252)
    at java.lang.Thread.join(Thread.java:1326)
    at org.apache.hadoop.hdfs.DataStreamer.closeResponder(DataStreamer.java:
973)
    at org.apache.hadoop.hdfs.DataStreamer.endBlock(DataStreamer.java:624)
    at org.apache.hadoop.hdfs.DataStreamer.run(DataStreamer.java:801)
18/09/30 06:20:46 INFO mapreduce.JobSubmitter: number of splits:9
18/09/30 06:20:47 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1538278733852_0007
18/09/30 06:20:47 INFO impl.YarnClientImpl: Submitted application
application_1538278733852_0007
18/09/30 06:20:47 INFO mapreduce.Job: The url to track the job: http://
ip-172-31-5-137.us-east-2.compute.internal:20888/proxy/
application_1538278733852_0007/
18/09/30 06:20:47 INFO mapreduce.Job: Running job: job_1538278733852_0007
18/09/30 06:21:13 INFO mapreduce.Job: Job job_1538278733852_0007 running in
uber mode : false
18/09/30 06:21:13 INFO mapreduce.Job: map 0% reduce 0%
18/09/30 06:21:32 INFO mapreduce.Job: map 15% reduce 0%
18/09/30 06:21:33 INFO mapreduce.Job: map 22% reduce 0%
18/09/30 06:21:36 INFO mapreduce.Job: map 33% reduce 0%
18/09/30 06:21:37 INFO mapreduce.Job: map 44% reduce 0%

```

18/09/30 06:21:38 INFO mapreduce.Job: map 67% reduce 0%
18/09/30 06:21:39 INFO mapreduce.Job: map 74% reduce 0%
18/09/30 06:21:42 INFO mapreduce.Job: map 81% reduce 0%
18/09/30 06:21:43 INFO mapreduce.Job: map 89% reduce 0%
18/09/30 06:21:47 INFO mapreduce.Job: map 93% reduce 0%
18/09/30 06:21:50 INFO mapreduce.Job: map 100% reduce 0%
18/09/30 06:21:52 INFO mapreduce.Job: map 100% reduce 25%
18/09/30 06:21:53 INFO mapreduce.Job: map 100% reduce 50%
18/09/30 06:21:55 INFO mapreduce.Job: map 100% reduce 100%
18/09/30 06:21:56 INFO mapreduce.Job: Job job_1538278733852_0007 completed successfully

18/09/30 06:21:56 INFO mapreduce.Job: Counters: 56

File System Counters

FILE: Number of bytes read=7954111
FILE: Number of bytes written=43348948
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=883
HDFS: Number of bytes written=0
HDFS: Number of read operations=9
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=81612666
S3: Number of bytes written=6175163
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0

Job Counters

Killed map tasks=1
Killed reduce tasks=1
Launched map tasks=9
Launched reduce tasks=4
Data-local map tasks=9
Total time spent by all maps in occupied slots (ms)=11008368
Total time spent by all reduces in occupied slots (ms)=5668032
Total time spent by all map tasks (ms)=229341
Total time spent by all reduce tasks (ms)=59042
Total vcore-milliseconds taken by all map tasks=229341
Total vcore-milliseconds taken by all reduce tasks=59042
Total megabyte-milliseconds taken by all map tasks=352267776
Total megabyte-milliseconds taken by all reduce tasks=181377024

Map-Reduce Framework

Map input records=1405935
Map output records=14203980
Map output bytes=202111353
Map output materialized bytes=33191208
Input split bytes=883
Combine input records=14203980
Combine output records=3820158
Reduce input groups=424462
Reduce shuffle bytes=33191208
Reduce input records=3820158
Reduce output records=424462
Spilled Records=7640316
Shuffled Maps =36
Failed Shuffles=0
Merged Map outputs=36

GC time elapsed (ms)=7339
CPU time spent (ms)=122440
Physical memory (bytes) snapshot=8967155712
Virtual memory (bytes) snapshot=48418066432
Total committed heap usage (bytes)=8195145728
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
Bytes Read=81612666
File Output Format Counters
Bytes Written=6175163