



```

Total vcore-milliseconds taken by all reduce tasks=7036
Total megabyte-milliseconds taken by all map tasks=288761856
Total megabyte-milliseconds taken by all reduce tasks=21614592
Map-Reduce Framework
  Map input records=1405946
  Map output records=6067351
  Map output bytes=47130252
  Map output materialized bytes=8332
  Input split bytes=985
  Combine input records=6067351
  Combine output records=820
  Reduce input groups=82
  Reduce shuffle bytes=8332
  Reduce input records=820
  Reduce output records=82
  Spilled Records=1640
  Shuffled Maps =10
  Failed Shuffles=0
  Merged Map outputs=10
  GC time elapsed (ms)=7774
  CPU time spent (ms)=45100
  Physical memory (bytes) snapshot=7122997248
  Virtual memory (bytes) snapshot=37727354880
  Total committed heap usage (bytes)=6552027136
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=81613300
File Output Format Counters
  Bytes Written=891
[hadoop@ip-172-31-5-137 wordcount]$ rm WordCount3.jar
[hadoop@ip-172-31-5-137 wordcount]$ hadoop jar WordCount3.jar com.wordcount3.WordCount s3://biblewordcount/ s3://biblewordcount/output2/
18/09/30 06:05:53 INFO client.RMProxy: Connecting to ResourceManager at ip-172-31-5-137.us-east-2.compute.internal/172.31.5.137:8032
18/09/30 06:05:55 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/09/30 06:05:55 INFO input.FileInputFormat: Total input files to process : 10
18/09/30 06:05:56 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library
18/09/30 06:05:56 INFO lzo.LzoCodec: Successfully loaded & initialized native-lzo library [hadoop-lzo rev 29dbe82f1ecfd8384b89f1a32e5e2d9e5dfd7724]
18/09/30 06:05:56 INFO mapreduce.JobSubmitter: number of splits:10
18/09/30 06:05:56 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1538278733852_0005
18/09/30 06:05:56 INFO impl.YarnClientImpl: Submitted application application_1538278733852_0005
18/09/30 06:05:56 INFO mapreduce.Job: The url to track the job: http://ip-172-31-5-137.us-east-2.compute.internal:20888/proxy/application_1538278733852_0005/
18/09/30 06:05:56 INFO mapreduce.Job: Running job: job_1538278733852_0005
18/09/30 06:06:07 INFO mapreduce.Job: Job job_1538278733852_0005 running in uber mode : false
18/09/30 06:06:07 INFO mapreduce.Job:  map 0% reduce 0%
18/09/30 06:06:24 INFO mapreduce.Job:  map 30% reduce 0%
18/09/30 06:06:25 INFO mapreduce.Job:  map 40% reduce 0%
18/09/30 06:06:29 INFO mapreduce.Job:  map 50% reduce 0%
18/09/30 06:06:31 INFO mapreduce.Job:  map 70% reduce 0%
18/09/30 06:06:32 INFO mapreduce.Job:  map 80% reduce 0%
18/09/30 06:06:33 INFO mapreduce.Job:  map 100% reduce 0%
18/09/30 06:06:34 INFO mapreduce.Job:  map 100% reduce 25%
18/09/30 06:06:37 INFO mapreduce.Job:  map 100% reduce 50%
18/09/30 06:06:38 INFO mapreduce.Job:  map 100% reduce 75%
18/09/30 06:06:39 INFO mapreduce.Job:  map 100% reduce 100%
18/09/30 06:06:39 INFO mapreduce.Job: Job job_1538278733852_0005 completed successfully
18/09/30 06:06:39 INFO mapreduce.Job: Counters: 55
File System Counters
  FILE: Number of bytes read=2023
  FILE: Number of bytes written=2398437
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=985
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=10
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=0
  S3: Number of bytes read=81613300
  S3: Number of bytes written=891
  S3: Number of read operations=0
  S3: Number of large read operations=0
  S3: Number of write operations=0
Job Counters
  Killed map tasks=1
  Launched map tasks=10
  Launched reduce tasks=4

```

```
Data-local map tasks=10
Total time spent by all maps in occupied slots (ms)=9160848
Total time spent by all reduces in occupied slots (ms)=3446592
Total time spent by all map tasks (ms)=190851
Total time spent by all reduce tasks (ms)=35902
Total vcore-milliseconds taken by all map tasks=190851
Total vcore-milliseconds taken by all reduce tasks=35902
Total megabyte-milliseconds taken by all map tasks=293147136
Total megabyte-milliseconds taken by all reduce tasks=110290944
Map-Reduce Framework
  Map input records=1405946
  Map output records=6067351
  Map output bytes=47130252
  Map output materialized bytes=9684
  Input split bytes=985
  Combine input records=6067351
  Combine output records=820
  Reduce input groups=82
  Reduce shuffle bytes=9684
  Reduce input records=820
  Reduce output records=82
  Spilled Records=1640
  Shuffled Maps =40
  Failed Shuffles=0
  Merged Map outputs=40
  GC time elapsed (ms)=7781
  CPU time spent (ms)=47820
  Physical memory (bytes) snapshot=8210546688
  Virtual memory (bytes) snapshot=51722678272
  Total committed heap usage (bytes)=7659323392
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=81613300
File Output Format Counters
  Bytes Written=891
[hadoop@ip-172-31-5-137 wordcount]$
```